PRODUCT TEAM BRIEFING · AUGUST 2026

Agentic Commerce

What it is, what's real, and what it means for what we build.

Agentic commerce is commerce where the buyer is software. A user states a goal — "trail running shoes, under $150, here by Friday" — and an AI agent handles some or all of discovery, comparison, cart building, checkout and post-purchase service.

Every layer of e-commerce assumed a human with eyes, a cursor, a cookie and a card. Agents break all four at once. That breakage is the whole industry — and the reason the roadmap implications are larger than they first appear.

THE SIX THINGS THAT ARE TRUE RIGHT NOW

1. **Discovery is real and large. Autonomous buying is not.** AI-referred traffic to US retailers grew **393% YoY in Q1 2026** and now converts *better* than conventional traffic. But only ~**12%** of consumers will let an agent complete a payment unsupervised. Today's money is in *agent-influenced* commerce, not agent-executed.
2. **The market leader publicly retreated.** OpenAI launched in-chat checkout (Sept 2025), added a 4% merchant fee (Jan 2026), then **killed buy-in-chat in March 2026** — fewer than 15 Shopify merchants ever went live, and Walmart measured in-chat checkout converting **~3× worse** than click-out to its own site. The model is now "**discover in the AI, transact on the merchant.**"
3. **Google and Shopify won the standards fight.** Their Universal Commerce Protocol (UCP) launched Jan 2026. By April, Amazon, Meta, Microsoft, Salesforce *and Stripe* — the rival protocol's own co-author — had joined its governance council.
4. **Payment rails are solved and largely free.** Visa, Mastercard, Amex, Stripe, PayPal and Coinbase have all shipped agent-payment infrastructure. The rails arrived ahead of the demand.
5. **The law arrived early and has teeth.** Amazon won a preliminary injunction against Perplexity's shopping agent in March 2026 under the Computer Fraud and Abuse Act. Undisclosed agents impersonating browsers are now a legal liability.
6. **Liability is unresolved and defaults to the merchant.** No US federal framework governs agent-initiated purchases. Merchants absorb the chargebacks.

The one-sentence version. The infrastructure for machine buyers is built and mostly free; the scarce assets are **consumer trust**, **the customer relationship**, and **the right to be recommended** — and that is what everyone is actually fighting over.

Source confidence. Much of what's published on this topic is vendor marketing. Claims below are tagged **VERIFIED** (spec, court record, earnings), **REPORTED** (credible press or analyst), or **SOFT** (vendor content — directional only, don't put it in a deck).

The model: five questions every agentic transaction must answer

Every protocol, vendor and lawsuit in this space is an answer to one of five questions. Learn the five and you can place any new announcement without memorising a vendor list.

```
┌─────────┐   ┌─────────┐   ┌─────────┐   ┌─────────┐   ┌─────────┐
│ 1. FIND │→  │2. INTENT│→  │3. TRUST │→  │ 4. PAY  │→  │ 5. FIX  │
│discovery│   │authority│   │identity │   │  rails  │   │recourse │
└─────────┘   └─────────┘   └─────────┘   └─────────┘   └─────────┘

 feeds, data   mandates,     agent auth,   scoped        disputes,
 structured    spend caps,   bot detect,   tokens,       returns,
 catalogs      human gates   signatures    networks      liability
   SOLVABLE      HARD          SOLVED        SOLVED        UNSOLVED
   TODAY       (the frontier)
```

QUESTION	WHAT IT MEANS	STATUS
1 · Find	How does an agent know your product exists? Agents don't browse, they read — structured feeds, schema markup, product identifiers.	**Where the value is today.** Mostly a data-quality problem, and therefore fixable.
2 · Intent	What did the human actually authorise? A human at checkout *is* the authorisation; an agent needs a portable, provable artefact of delegated authority.	The genuinely hard problem and the real frontier.
3 · Trust	Is this agent legitimate? Answered cryptographically — agents sign their requests, merchants verify against a key directory.	Technically solved; adoption spreading.
4 · Pay	How does money move without exposing a card? Every rail converged on the same primitive: a credential scoped to one merchant, one amount, minutes of validity.	Solved. Multiple competing implementations.
5 · Fix	When it goes wrong, who pays? Dispute frameworks assume a human decided.	**Unsolved.** This is what caps autonomy.

Why this matters for planning: roughly 80% of announcements in this space address Q1 or Q3, because Q2 and Q5 require someone to accept liability and nobody wants to. When you see a launch, ask which question it answers. If it's Q1/Q3, it's incremental. If it credibly addresses Q2 or Q5, it's significant.

The protocol landscape

Two rival commerce standards, plus a supporting cast that is *complementary*, not competitive — a distinction most coverage gets wrong.

	ACP	UCP
Backers	OpenAI, Stripe, Meta	Google, Shopify + 8-member council
Launched	Sept 2025	Jan 2026
Design goal	Complete a checkout inside the AI surface	Negotiate any commerce interaction across the full lifecycle
Scope	Product feed, checkout session, delegated payment, order webhooks	Discovery, catalog, cart, checkout, post-purchase, loyalty, fulfillment
Human handoff	Not a first-class concept	First-class — explicit "needs a human" state and handoff
Flagship surface	ChatGPT in-chat checkout — **retired Mar 2026**	Google AI Mode, Gemini, Universal Cart, Microsoft Copilot Checkout
Merchant cost	OpenAI charged 4% on checkout sales	No additional platform fee on Google/Microsoft surfaces — *for now*
Status	Alive, but a spec looking for a surface	**The default.**

ALSO KNOW	OWNER	ANSWERS	WHAT IT DOES
MCP	Anthropic	—	Standard way for an agent to call external tools and data. The substrate everything else runs on. Cheapest way to become agent-addressable.
A2A	Google	—	Agent-to-agent negotiation. Enables buyer-agent ↔ merchant-agent interaction.
AP2	Google + 60	Q2, Q4	Chain of signed "mandates" (intent → cart → payment) proving what the user authorised. The conceptual answer to delegated authority.
Visa TAP	Visa + Cloudflare	Q3	Cryptographic agent identity in signed HTTP headers, verified against a key directory.
x402	Coinbase	Q4	Stablecoin micropayments over HTTP. Real technology; real demand still thin.

The strategic read: UCP won **governance**, which is more durable than early transaction volume. Note what Stripe did — joined UCP's council without abandoning ACP. That's infrastructure behaviour: stay neutral, be everywhere, tax the flow regardless of who wins. Larger merchants increasingly run **both**.

Reality check: what's working and what isn't

Both sets of numbers below are correct. They measure different stages, and conflating them is the single most common error in this space.

```
INFLUENCE  ████████████████████████████████  large · growing fast · converts well
EXECUTION  ███                               small · contested · converts worse

The industry markets EXECUTION.   The money is in INFLUENCE.
```

WORKING — STRONG EVIDENCE	NOT WORKING — EQUALLY STRONG EVIDENCE
• AI-referred traffic +393% YoY (Q1 2026); +693% over the 2025 holidays REPORTED • That traffic engages 12% more, spends 48% longer, and generates **37% higher revenue per visit** • Conversion swung from 38% *worse* than direct (Mar 2025) to 42% *better* (Mar 2026) • Salesforce: AI *influenced* >20% of global online retail sales, 2025 holidays • Amazon's own assistant (Rufus) drove ~**$12B** incremental annualised sales in 2025 VERIFIED	• Only **12%** of consumers accept autonomous decisions at the payment stage; 32% pre-payment; 74% for routine tasks (Accenture, 25,590 consumers, 16 countries) • Forrester (mid-2026): "very few consumers allow agents to complete purchases without direct oversight, and even fewer do so regularly" • Retail sites still aren't machine-readable — Adobe's own finding • Gartner projects a large share of agentic commerce projects cancelled by 2027 REPORTED • In-surface checkout has underperformed click-out wherever it's been measured

CLAIMS TO TREAT AS UNPROVEN

"Agents will autonomously buy groceries" · "$1T market by 2030" (assumes autonomy that doesn't exist) · "Integrate checkout now or die" (the leader just reversed) · "Crypto rails are winning agent payments" (transaction counts ≠ demand) · "One protocol will win outright."

HOW TO READ ANY ADOPTION STATISTIC

Survey results swing 60 points on wording. "*Would you let AI shop for you?*" → ~70% yes. "*Would you let it complete payment unsupervised?*" → ~12%. Before quoting a number, ask: **at which funnel stage was this measured, and against what baseline?**

What this means for what we build

Ordered by return, given the evidence above.

1 — Product data is the highest-ROI work, and it isn't close

Agents filter on constraints. Every missing attribute is a filter you silently fail. Adobe's finding is that AI traffic is surging while most retail sites remain not machine-readable — that gap is the opportunity.

- **Feed quality first.** Google Merchant Center is now the primary product-data source for both Gemini and ChatGPT REPORTED — one artefact, multiple surfaces.
- **Structured data:** schema.org Product markup, GS1 standards, correct identifiers, explicit variant modelling.
- **Freshness matters as a ranking input.** Stale inventory or pricing means the agent recommends you, the purchase fails, and the surface deprioritises you.
- **Machine-readable policies** — shipping, returns, cancellation, delivery promise. Agents weigh these explicitly.

2 — Be agent-addressable before you're agent-integrated

- **Cheapest meaningful step:** expose search, inventory and cart via an MCP server. Low cost, works with any capable agent, commits you to no protocol.
- **Protocol choice:** UCP is the default. Add ACP only if a surface you specifically care about requires it.
- **Revisit bot policy.** Blanket-blocking automation now costs real demand. Allow *signed, identified* agents; block unsigned ones.

3 — Instrument for the new dispute regime

Capture per agentic order: agent identity, the authorisation artefact, the action log, the exact cart at approval, delivery proof. Merchants who can produce this evidence pack win disputes that competitors lose. Nobody is doing this well yet.

4 — Measure the things that now matter

METRIC	WHAT IT TELLS YOU
Found rate	Share of tested prompts where an AI surfaces your brand at all — top of the new funnel
Recommendation rate	Share where you're actually recommended. Found ≠ recommended.
AI share of voice	Your share of AI answers vs. competitors on a defined query set
Agent conversion rate	Agent-led sessions ending in purchase — compare against your own click-out baseline. That comparison is what killed in-chat checkout.
Feed error rate	SKUs with missing or invalid data. A direct, fixable revenue leak.

> **THE POSTURE THAT MATCHES THE EVIDENCE**
> **Optimise hard for discovery. Be deliberate about ceding checkout.** Discovery traffic is large, growing and converting well. Checkout inside someone else's surface costs the customer relationship, the login, the loyalty hook and the data — and so far converts worse. Build machine advantage in data and content; experiment with assistive AI on our own surfaces first.

Risk: why autonomy has a ceiling

Three constraints, all of which bound how far this can go regardless of how good the models get.

PROMPT INJECTION IS STRUCTURALLY UNSOLVED

An LLM can't reliably distinguish instructions from data — both arrive as text. In commerce, that means a product description, review or page element can carry hidden instructions that redirect an agent. OWASP research finds this still drives most agentic AI security failures in production, and in May 2026 the Five Eyes agencies issued joint guidance naming it a core attack path with no single sufficient safeguard. **You cannot responsibly give unsupervised spending authority to a system that can be talked out of its instructions by a product page.**

LIABILITY DEFAULTS TO THE MERCHANT

Dispute frameworks assume a human decided. When an agent purchases autonomously under authority granted days earlier, that assumption collapses and there's no field for it. The CFPB's Jan 2026 advisory held that consumer dispute rights *survive* delegation to an agent REPORTED. Amex's commitment to cover erroneous purchases by registered agents is the notable early exception — and a signal of where this has to go.

AGENT ACCESS IS NOT A RIGHT

Amazon v. Perplexity established that a platform can lawfully exclude third-party agents, and that **concealment is the aggravating fact** — the legal problem wasn't automation, it was impersonating a human browser. Consent-first is now the defensible posture, and disclosure is the price of admission.

> **The commercial consequence:** these three constraints are why the "agents buy everything autonomously" scenario keeps slipping. They're not engineering backlog items — they're structural. Plan for **assisted** commerce, and treat full autonomy as an option you buy later, not a near-term default.

Economics: who captures the margin

The organising question in this whole space is **who captures the margin when the buyer is a machine.** Two structural shifts matter for planning.

ATTRIBUTION IS BREAKING

The ~$13B US affiliate industry runs on cookies, clicks and last-touch attribution. Agents follow no affiliate links, carry no cookies and travel no referral paths. The classic ~14-click journey compresses to 1–2 interactions. Replacements are emerging — platform transaction fees, outcome-based commissions tied to incremental sales, crediting the data source that informed the decision — but none is standard yet.

THE FINDING THAT SHOULD CONCERN ANYONE WITH A RETAIL MEDIA PLAN

Research reported out of Columbia/Yale found **sponsored tags reduce an agent's probability of selecting a product by 11–21%** across models REPORTED. Agents actively *penalise* paid placement, because they're optimising the user's stated goal rather than platform yield. If that holds, the placement-based ad model doesn't port to agentic surfaces — monetisation has to become performance-based.

PLAYER	POSITION	VERDICT
Agent surfaces Google, OpenAI, MSFT, Meta	Own demand and ranking. Google also owns the product-data substrate and the winning protocol.	Strongest hand
Card networks	Moved earliest; own identity, tokenization and risk pricing.	Defended successfully
PSPs Stripe, Adyen, PayPal	Neutral infrastructure across every protocol.	Win either way
Commerce platforms Shopify	Co-authored UCP; distribution as leverage.	Kingmaker
Large retailers	Have the data to say no. Walmart's conversion finding reshaped the market.	Strong, if they hold the customer
Long-tail merchants	Bear integration cost, gain reach, lose the relationship and the data.	Weakest position

On market sizing: the trillion-dollar projections assume autonomy levels consumer research says don't exist. The more defensible near-term figure is eMarketer's ~**$20.9B** of 2026 US retail spending attributed to AI platforms — roughly 4× 2025 REPORTED.

What to watch — and what to treat as stable

This briefing is a snapshot of a field that moves quarterly. The distinction below is what keeps it useful rather than misleading.

TREAT AS STABLE	EXPECT TO CHANGE WITHIN MONTHS
• Delegated authority must be provable before anyone lets software spend • Liability determines autonomy • Prompt injection caps unsupervised spending • Whoever controls demand captures the margin • Machines read structured data; ambiguity a human infers is invisible to an agent • The five-question model • Influence vs. execution as the distinction that resolves contradictory stats	• Every specific number in this document • Which protocol leads, and its version • Fee structures — today's "free" surfaces will monetise • Which surfaces support checkout natively • Tool and vendor recommendations • Who sits on which governance body

LEADING INDICATORS WORTH TRACKING

- **Google's Universal Cart conversion vs. click-out.** Best-resourced attempt at in-surface checkout, from the party that owns the product data. If it beats click-out, in-surface checkout is viable and OpenAI simply executed badly. If not, "discover in AI, buy on merchant" is permanent.
- **Whether a card network writes agentic dispute rules.** Liability is the gate on autonomy; formal rules are the unlock.
- **Real human-not-present volume.** The protocols support agents buying against pre-set rules. Nobody uses it at scale. That volume is the only true proof this is more than conversational discovery.
- **Whether monetisation goes performance-based** rather than placement-based.
- **Amazon's posture** — it blocks third-party agents, won an injunction, runs its own assistant, and joined UCP's council. Which way it opens matters.

Vocabulary

Enough to be fluent in a meeting.

- **ACP** — Agentic Commerce Protocol. OpenAI/Stripe/Meta. Checkout + delegated payment.
- **UCP** — Universal Commerce Protocol. Google/Shopify. Full lifecycle. The default.
- **AP2** — Agent Payments Protocol. Signed mandate chain proving authorisation.
- **MCP** — Model Context Protocol. How agents call tools and data.
- **A2A** — Agent-to-agent negotiation.
- **Mandate** — cryptographically signed record of what a user authorised.
- **Human-present / not-present** — user approves the actual cart, vs. agent executes later against pre-signed rules.
- **Scoped token** — payment credential locked to one merchant, one amount, minutes of validity.
- **Merchant of record** — who legally sells. Stays the merchant, not the agent.
- **Product feed** — structured file of what you sell. The single highest-leverage asset.
- **GEO / AEO / ACO** — optimising to be surfaced by, cited by, and transactable with AI.
- **Found rate / recommendation rate** — surfaced at all, vs. actually recommended.
- **Prompt injection** — hidden text hijacking an agent's instructions.
- **Evidence pack** — the dispute artefacts: authorisation, agent identity, action log, cart, delivery proof.

If you want to go deeper

- **Shopify Engineering, "Building UCP"** — shopify.engineering/ucp. The best technical document in the field.
- **Stripe ACP docs** — docs.stripe.com/agentic-commerce/acp
- **Adobe AI traffic reports** (quarterly, free) — the most-cited hard numbers.
- **Forrester, "The State Of Agentic Commerce In Mid-2026"** — the best sceptical counterweight.
- **Cloudflare, "Securing agentic commerce"** — clearest explanation of agent identity.
- **Google AP2 announcement** — the mandate model, explained by its authors.

> **The takeaway for our roadmap:** treat agentic commerce as a **discovery channel** that is real, large and growing **—** and as an **execution channel** that is early, contested, and not yet worth ceding the customer relationship for. The work that pays today is making our data machine-readable. The work that pays later is being ready to transact when liability, trust and conversion catch up.